

Big Data

Spark RDD

Spark RDD

❖ RDD (Resilient Distributed Datasets)

- RDD General Characteristics
 - Sequence of deterministic Transformations & Actions are conducted on a
 - Stored dataset
 - Another RDD
 - Group of RDDs
 - RDD = Resilient + Distributed + Datasets

Spark RDD

❖ RDD (Resilient Distributed Datasets)

- Resilient
 - Fault-tolerant collection of data
 - RDDs can recover from node failures automatically
 - Immutable distributed collection of objects/data

Spark RDD

❖ RDD (Resilient Distributed Datasets)

- Distributed
 - Parallel processing of distributed data sets
 - Each logical partition of the dataset can be computed on different nodes of the cluster
 - Transformations are Lazy during DAG (Directed Acyclic Graph) setup

Spark RDD

❖ RDD (Resilient Distributed Datasets)

- Distributed
 - Persisting (`persist()`) and Caching (`cache()`) operations enable secure, fast, and very easy to program dataset distribution and sharing among multiple nodes in the cluster

Spark RDD

❖ RDD (Resilient Distributed Datasets)

- Dataset
 - RDD divides each dataset into logical partitions
 - Read-only based partitioned collection of data
 - All types of Python, Java, and Scala objects (including user-defined classes) can be processed in RDDs

Spark RDD

❖ RDD Immutable Data

- Immutable data means that it can be recreated again if needed
- Immutable data comes from a set of deterministic Transformations applied to an input dataset
- Immutability makes saving, copying, and sharing easy

Spark RDD

❖ RDD Immutable Data

- Immutability helps solve simultaneous multiple thread update problems
- Immutable data can be safely shared and distributed across processes and nodes

Spark RDD

❖ RDD Persistence

- Persisting (or Caching) a dataset in memory enhances data processing speed across multiple operations
- Persisted RDD makes a node(s) store its Transformed data in its partition memory to quickly reuse it for other Actions processed on that dataset

Spark RDD

❖ RDD Persistence

- Using a DAG (Directed Acyclic Graph) and having Lazy Transformations helps to schedule optimized RDD Persisting for multiple Actions to be executed more efficiently

Spark RDD

❖ RDD Components

- Partitions
 - Logically divided chunk of the RDD or a larger dataset (may be a distributed dataset)
 - All RDDs are divided into Partitions
 - RDD's main unit of data processing
 - RDD automatically partitions datasets

Spark RDD

❖ RDD Components

- Partitions
 - Programmer can control the partition number and size to more efficiently support the application
 - Spark RDD default partition size for HDFS is 64 MB (HDFS standard block size)
 - Spark RDD maximum partition size for HDFS is 2 GB

Spark RDD

❖ RDD Components

- **Metadata & Data Type**
 - Data and object type considered in the RDD partitioning and data placement
 - Structured data, Semi-structured data, and Unstructured data
- **Task & Transformation**
 - Function applied to process the data partition
 - Parent RDD → Function → Child RDD
 - Parent : Child → 1:1, N:1, 1:M, N:M

Spark RDD

❖ RDD Types

- | | | |
|------------------------------------------------------------------------------------|----------------------|-----------------------------------|
| ▪ HadoopRDD
Each HDFS block is a partition | ▪ ShuffledRDD | ▪ VertexRDD |
| ▪ FilteredRDD
Child RDD partitions are same as the parent RDD partitions | ▪ UnionRDD | ▪ EdgeRDD |
| ▪ JoinedRDD
One partition per Reduce task | ▪ PythonRDD | ▪ CassandraRDD |
| ▪ MappedRDD | ▪ DoubleRDD | ▪ GeoRDD |
| ▪ PairRDD | ▪ JdbcRDD | ▪ EsSpark
ElasticSearch |
| | ▪ JsonRDD | ▪ etc. |
| | ▪ SchemaRDD | |

Spark RDD

❖ SchemaRDD

- RDD that has an associated schema
- Has all standard RDD functions and also can be used in relational queries

Spark RDD

❖ SchemaRDD creation methods

- Loading data in from external sources to SchemaRDD form
- Methods of converting a standard RDD into a SchemaRDD
- Importing a SQLContext
- Apply the createSchemaRDD function on SQLContext

Spark RDD

❖ SchemaRDD

- SchemaRDD can be registered as a table in the SQLContext it was imported from
- SchemaRDD registered as a table can be used in SQL statements

Big Data
References

References

- Holden Karau, Andy Konwinski, Patrick Wendell, and Matei Zaharia, *Learning Spark: Lightning-Fast Big Data Analysis*. 1st Edition. O'Reilly, 2015.
- Sameer Farooqui, Databricks, **Advanced Apache Spark Training**, Devops Advanced Class, Spark Summit East 2015, <http://slideshare.net/databricks>, www.linkedin.com/in/blueplastic, March 2015.
- Apache Spark documents (all documents and tutorials were used)
 - <http://spark.apache.org/docs/latest/rdd-programming-guide.html>
 - <http://spark.apache.org/docs/latest/rdd-programming-guide.html#working-with-key-value-pairs>
 - <https://spark.apache.org/docs/2.2.0/rdd-programming-guide.html#rdd-persistence>
- Wikipedia, www.wikipedia.org
- Stackoverflow, <https://stackoverflow.com/questions>
- Bernard Marr, "Spark Or Hadoop -- Which Is The Best Big Data Framework?," Forbes, Tech, June 22, 2015.
- Quick introduction to Apache Spark, <https://www.youtube.com/watch?v=TgiBvKcGL24>
- Wide vs Narrow Dependencies, <https://github.com/rohgar/scala-spark-4/wiki/Wide-vs-Narrow-Dependencies>

References

- Partitions and Partitioning, <https://jaceklaskowski.gitbooks.io/mastering-apache-spark/spark-rdd-partitions.html>
- Neo4j, "From Relational to Neo4j," <https://neo4j.com/developer/graph-db-vs-rdbms/> (last accessed Jan. 1, 2018).

Image Sources

- By Robivy64 at English Wikipedia [Public domain], via Wikimedia Commons
- Teravolt at English Wikipedia [CC BY 3.0 (<http://creativecommons.org/licenses/by/3.0/>)], via Wikimedia Commons
- By Konradr (Own work) [GFDL (<http://www.gnu.org/copyleft/fdl.html>) or CC-BY-SA-3.0 (<http://creativecommons.org/licenses/by-sa/3.0/>)], via Wikimedia Commons