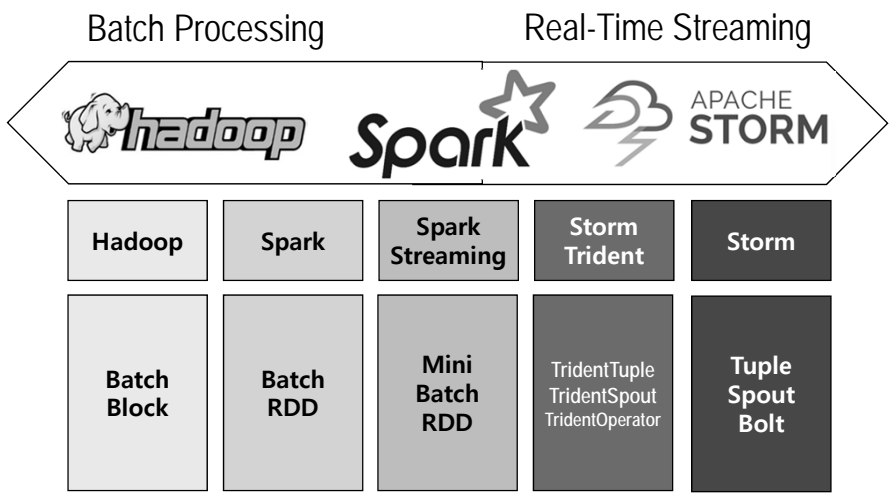


Big Data

Spark Streaming

Spark Streaming



Spark Streaming

❖ Spark Streaming Characteristics

- Extension to the Spark Core API
- Live data streams can be processed
- Fault-tolerant and scalable
- High throughput (near) real-time data processing → 0.5 s or longer
- Streaming data input from HDFS, Kafka, Flume, TCP sockets, Kinesis, etc.

Spark Streaming

❖ Spark Streaming Characteristics

- Stream processing high-level functions
 - Map, Reduce, Join, Window, etc.
- Processed Output saved on Filesystems, Databases, and Live Dashboards
- Spark ML (Machine Learning) functions and GraphX graph processing algorithms are fully applicable to streaming data

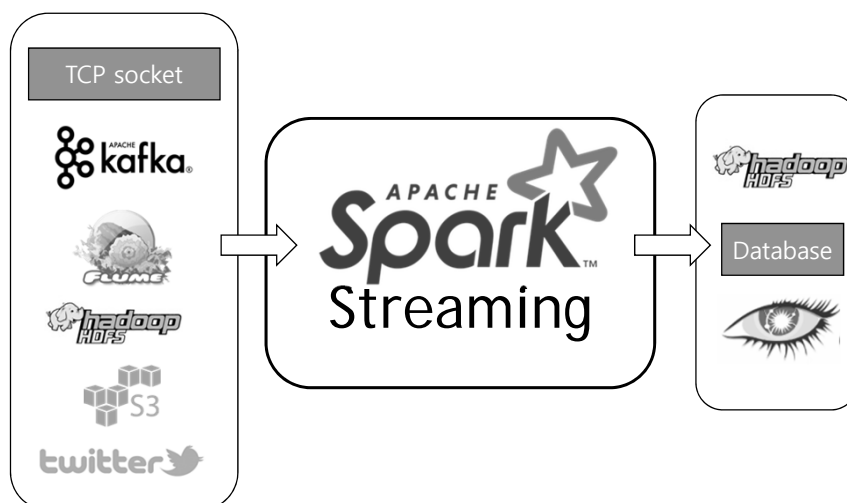
Spark Streaming

❖ Spark Streaming Characteristics

- Spark uses (size controllable) micro-batch processing of data for real-time analysis
 - Hadoop uses batch processing of data, which is time consuming to obtain results
- Spark uses RDD to arrange data and recover from failures

Spark Streaming

❖ Spark Streaming Input & Output



Spark Streaming

❖ Streaming Receiver Types

- Basic
 - File systems
 - Socket connectors
 - Akka Actors
 - Sources directly available in Streaming Context API

Spark Streaming

❖ Streaming Receiver Types

- Custom
 - Requires implementing an user-defined receiver
- Advanced
 - Requires linking with systems that have extra dependencies
 - Kafka, Flume, Twitter

Spark Streaming

❖ Spark Streaming process

1. Live input data stream received
2. Input data stream is divided into Mini-Batches called a DStream (Discretized Stream), which is saved as a small RDD every mini-batch period
3. Spark Stream engine cores process the mini-batches and generate a final output stream of mini-batches

Spark Streaming

❖ DStream

- DStream (Discretized Stream) is a continuous stream of data with high-level abstraction
- DStreams are created from input data stream sources (Kafka, Flume, Kinesis, etc.) or high-level processing operations on other DStreams

Spark Streaming


❖ DStream

- DStreams are represented as a sequence of small RDDs
- Mini-Batch size is 0.5 s or longer
- RDD is processed through the DAG
- Processing latency (through the DAG) has to be smaller than the mini-batch period

Spark Streaming

❖ Window Operations

- Window Length
 - Number of blocks (partitions) to conduct a RDD DAG process together



- Sliding Interval
 - Number of blocks (partitions) to slide the Window after a RDD process is conducted

Spark Streaming


❖ Transformations for Windows

- Key parameters: `windowLength`, `slideInterval`
- `window()`
- `countByWindow()`
- `reduceByWindow()`
- `reduceByKeyAndWindow()`
- `countByValueAndWindow()`
- etc.

Spark Streaming

❖ Window Operations

- If Window Length > Sliding Interval, then
Overlap Blocks = (Window Length – Sliding Interval)
will exist for each RDD process



- Overlap Blocks help to analyze correlation (dependency) of sequential blocks of the streamed data

Spark Streaming

❖ Spark Streaming Examples

- IPTV or Web Page Live statistics
 - Channel or Page view of clicks
 - Use Kafka for buffering
 - Spark Streaming for processing
 - Draw a Heap Map of the current Channel or Page view clicks



Spark Streaming

❖ Spark Streaming Examples

- Sales Product Type Monitoring

1. Online Sales

- Read through Kafka




2. Department Store Sales

- Read through Flume

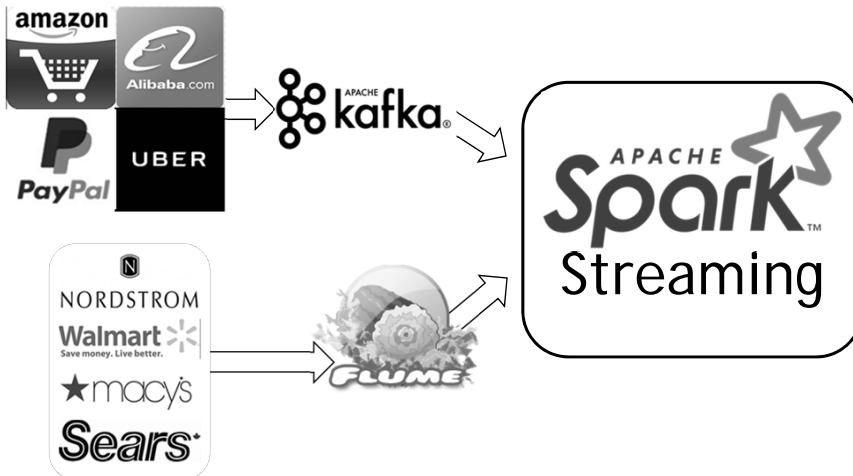


- Join 2 live data streams into Spark Streaming



Spark Streaming

❖ Streaming Receiver Types



Spark Streaming

❖ One Stream Input (e.g., from)

1. One Task slot in the Executor will serve as a Receiver (thread) to receive the live streaming data into a Block (Partition) of the RDD on the node

Spark Streaming

❖ One Stream Input (e.g., from kafka.)

2. Receiver will also make a copy of this Partition to another node (e.g., replication factor 2)
3. DAG Transformations are executed on the new RDD

Spark Streaming


❖ Another Stream Added (e.g., from)

1. On a different node, assign one Task slot in the Executor to serve as a Receiver (thread) to receive the live streaming data into a Block (Partition) of the RDD on the node
2. Receiver will also make a copy of this Partition to another node (e.g., replication factor 2)

Spark Streaming

❖ Another Stream Added (e.g., from )

3. DAG Transformations are executed on the new RDD
4. Union can be used to unify the two RDDs into one RDD



Big Data
References

References

- Holden Karau, Andy Konwinski, Patrick Wendell, and Matei Zaharia, *Learning Spark: Lightning-Fast Big Data Analysis*. 1st Edition. O'Reilly, 2015.
- Sameer Farooqui, Databricks, **Advanced Apache Spark Training**, Devops Advanced Class, Spark Summit East 2015, <http://slideshare.net/databricks>, www.linkedin.com/in/blueplastic, March 2015.
- Apache Spark documents (all documents and tutorials were used)
 - <http://spark.apache.org/docs/latest/rdd-programming-guide.html>
 - <http://spark.apache.org/docs/latest/rdd-programming-guide.html#working-with-key-value-pairs>
 - <https://spark.apache.org/docs/2.2.0/rdd-programming-guide.html#rdd-persistence>
- Wikipedia, www.wikipedia.org
- Stackoverflow, <https://stackoverflow.com/questions>
- Bernard Marr, "Spark Or Hadoop -- Which Is The Best Big Data Framework?," Forbes, Tech, June 22, 2015.
- Quick introduction to Apache Spark, <https://www.youtube.com/watch?v=TgiBvKcGL24>
- Wide vs Narrow Dependencies, <https://github.com/rohgar/scala-spark-4/wiki/Wide-vs-Narrow-Dependencies>

References

- Partitions and Partitioning, <https://jaceklaskowski.gitbooks.io/mastering-apache-spark/spark-rdd-partitions.html>
- Neo4j, "From Relational to Neo4j," <https://neo4j.com/developer/graph-db-vs-rdbms/> (last accessed Jan. 1, 2018).

Image Sources

- By Robivy64 at English Wikipedia [Public domain], via Wikimedia Commons
- Teravolt at English Wikipedia [CC BY 3.0 (<http://creativecommons.org/licenses/by/3.0/>)], via Wikimedia Commons
- By Konradr (Own work) [GFDL (<http://www.gnu.org/copyleft/fdl.html>) or CC-BY-SA-3.0 (<http://creativecommons.org/licenses/by-sa/3.0/>)], via Wikimedia Commons

Course Title

Big Data Emerging Technologies

❖ Modules

1. Big Data Rankings & Products
2. Big Data & Hadoop
3. Spark
4. Spark ML & Streaming
5. Storm
6. IBM SPSS Statistics Project